

# ML\_Assignment\_2\_Task\_1\_SalehAlkhalifa

October 9, 2024

## 1 Assignment 2 - Task 1

Table of Contents:

1. Import Libraries and Data
2. Exploratory Data Analysis
3. Data Preprocessing
4. OLS Regression
5. Ridge Regression
6. Lasso Regression
7. Comparison Analysis

### 1.0.1 1. Import Libraries and Data:

```
[309]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
```

```
[310]: df = pd.read_csv("HousingData.csv")
```

### 1.0.2 2. Exploratory Data Analysis:

```
[311]: df.head()
```

```
[311]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2

```
[312]: # I cannot remember which feature is which, so I am renaming them for
↳additional context
column_mapping = {
    "CRIM": "crime_rate",
    "ZN": "residential_land",
    "INDUS": "business_land",
    "CHAS": "by_river",
    "NOX": "nox_concentration",
    "RM": "avg_rooms",
    "AGE": "old_homes",
    "DIS": "job_distance",
    "RAD": "highway_access",
    "TAX": "property_tax",
    "PTRATIO": "student_teacher_ratio",
    "B": "black_population",
    "LSTAT": "low_status_pct",
    "MEDV": "median_value"
}

df = df.rename(columns=column_mapping)
```

```
[313]: # Getting a sense of the rows and columns
df.shape
```

```
[313]: (506, 14)
```

```
[314]: # Getting a sense of the data within the columns
df.describe()
```

```
[314]:
```

	crime_rate	residential_land	business_land	by_river	\
count	486.000000	486.000000	486.000000	486.000000	
mean	3.611874	11.211934	11.083992	0.069959	
std	8.720192	23.388876	6.835896	0.255340	
min	0.006320	0.000000	0.460000	0.000000	
25%	0.081900	0.000000	5.190000	0.000000	
50%	0.253715	0.000000	9.690000	0.000000	
75%	3.560263	12.500000	18.100000	0.000000	
max	88.976200	100.000000	27.740000	1.000000	

	nox_concentration	avg_rooms	old_homes	job_distance	\
count	506.000000	506.000000	486.000000	506.000000	
mean	0.554695	6.284634	68.518519	3.795043	
std	0.115878	0.702617	27.999513	2.105710	
min	0.385000	3.561000	2.900000	1.129600	
25%	0.449000	5.885500	45.175000	2.100175	
50%	0.538000	6.208500	76.800000	3.207450	
75%	0.624000	6.623500	93.975000	5.188425	

max	0.871000	8.780000	100.000000	12.126500
-----	----------	----------	------------	-----------

	highway_access	property_tax	student_teacher_ratio	black_population \
count	506.000000	506.000000	506.000000	506.000000
mean	9.549407	408.237154	18.455534	356.674032
std	8.707259	168.537116	2.164946	91.294864
min	1.000000	187.000000	12.600000	0.320000
25%	4.000000	279.000000	17.400000	375.377500
50%	5.000000	330.000000	19.050000	391.440000
75%	24.000000	666.000000	20.200000	396.225000
max	24.000000	711.000000	22.000000	396.900000

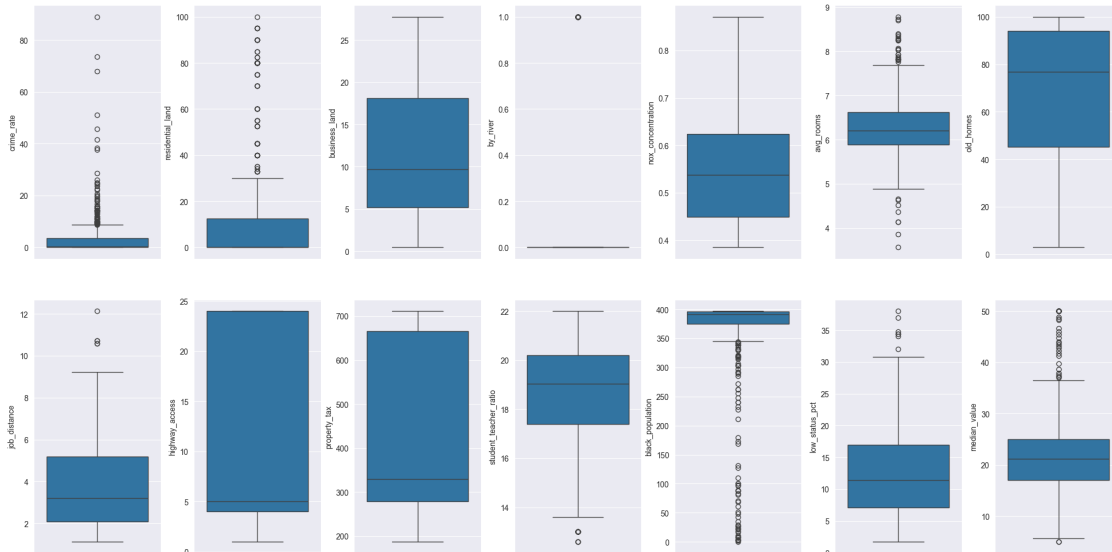
	low_status_pct	median_value
count	486.000000	506.000000
mean	12.715432	22.532806
std	7.155871	9.197104
min	1.730000	5.000000
25%	7.125000	17.025000
50%	11.430000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
[315]: # Understanding the data types
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   crime_rate             486 non-null    float64
1   residential_land        486 non-null    float64
2   business_land           486 non-null    float64
3   by_river                486 non-null    float64
4   nox_concentration       506 non-null    float64
5   avg_rooms               506 non-null    float64
6   old_homes               486 non-null    float64
7   job_distance            506 non-null    float64
8   highway_access          506 non-null    int64
9   property_tax            506 non-null    int64
10  student_teacher_ratio   506 non-null    float64
11  black_population        506 non-null    float64
12  low_status_pct          486 non-null    float64
13  median_value            506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```


```
[316]: # Plotting the distributions of the features
```

```
fig, axis = plt.subplots(ncols=7, nrows=2, figsize=(20, 10))
idx = 0
axis = axis.flatten()
for key, val in df.items():
    sns.boxplot(y=key, data=df, ax=axis[idx])
    idx += 1
plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=5.0)
```



```
[317]: plt.figure(figsize=(20, 10))
sns.heatmap(df.corr().abs(), annot=True)
```

```
[317]: <Axes: >
```



```
[318]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler, MinMaxScaler, PowerTransformer
from sklearn.decomposition import PCA
import numpy as np
from sklearn.model_selection import train_test_split

class BostonHousingEDA:
    def __init__(self, data):
        """
        Initializes the EDA class with the data for Boston Housing
        """
        # Set attributes
        self.data = data
        self.scaled_data = None
        self.cleaned_data = None
        self.X = None
        self.y = None
        self.column_mapping = {
            "CRIM": "crime_rate",
            "ZN": "residential_land",
            "INDUS": "business_land",
            "CHAS": "by_river",
            "NOX": "nox_concentration",
```

```

        "RM": "avg_rooms",
        "AGE": "old_homes",
        "DIS": "job_distance",
        "RAD": "highway_access",
        "TAX": "property_tax",
        "PTRATIO": "student_teacher_ratio",
        "B": "black_population",
        "LSTAT": "low_status_pct",
        "MEDV": "median_value"
    }
    self.data = self.data.rename(columns=self.column_mapping)

    # Define X and y
    self.X = self.data.drop(columns=['median_value'])
    self.y = self.data['median_value']

def handle_missing_values(self, strategy="mean"):
    """
    Handles missing values in the dataset with three methods
    """
    if strategy == "mean":
        self.data.fillna(self.data.mean(), inplace=True)
    elif strategy == "median":
        self.data.fillna(self.data.median(), inplace=True)
    elif strategy == "drop":
        self.data.dropna(inplace=True)
    else:
        raise ValueError("Error. Select value")

def visualize_data_distribution(self):
    """
    Visualize the distribution of all features in the dataset to get a
    ↪sense of shape
    """
    # Create the fig
    plt.figure(figsize=(20, 15))

    # Iterate over the columns and plot each
    for i, col in enumerate(self.data.columns, 1):
        # Create a subplot
        plt.subplot(4, 4, i)
        # Plot the histplot
        sns.histplot(self.data[col], kde=True, bins=30)
        plt.title(f'Distribution of {col}')
    plt.tight_layout()
    plt.show()

```

```

def visualize_correlation_matrix(self):
    """
    Plot the correlation matrix
    """
    # Get correlation
    corr_matrix = self.data.corr()
    # Create fig
    plt.figure(figsize=(12, 8))
    # Create heatmap
    sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
    plt.title('Correlation Matrix')
    plt.show()

def handle_multicollinearity(self, threshold=0.75):
    """
    Identify and remove highly correlated features based on a given
    ↪ threshold.
    """
    # add target y
    self.y = self.data['median_value']

    # create a copy of the features
    self.X = self.data.drop(columns=['median_value'])

    # calc the correlation matrix of the features
    corr_matrix = self.X.corr().abs()
    upper_tri = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).
    ↪ astype(bool))

    # drop with high corr
    to_drop = [column for column in upper_tri.columns if
    ↪ any(upper_tri[column] > threshold)]
    self.cleaned_data = self.X.drop(columns=to_drop)

    # update self.X with cleaned data
    self.X = self.cleaned_data

    print(f"The columns dropped due to MCL: {to_drop}")
    print(f"Remaining columns: {self.cleaned_data.columns.tolist()}")

    return self.cleaned_data, self.y

def scale_data(self, method='standard'):
    """
    Scale the data using the selected method and return scaled X and y

```

```

"""

# If/else with three options to compare for assignemtn
if method == 'standard':
    scaler = StandardScaler()
elif method == 'minmax':
    scaler = MinMaxScaler()
elif method == 'power':
    scaler = PowerTransformer()
else:
    raise ValueError("Error. Pick one")

# Use cleaned data
data_to_scale = self.X if self.X is not None else self.data.
↳drop(columns=['median_value'])
    self.scaled_data = pd.DataFrame(scaler.fit_transform(data_to_scale),
↳columns=data_to_scale.columns)

# return X and y for splitting
return self.scaled_data, self.y

def pca_visualization(self, n_components=2):
    """
    Apply the PCA algo to visualize the data
    """

    # Check for scaled data
    if self.scaled_data is None:
        print("Please scale the data first.")
        return

    # Apply PCA
    pca = PCA(n_components=n_components)
    pca_components = pca.fit_transform(self.scaled_data)

    # Plot results
    plt.figure(figsize=(8, 6))
    plt.scatter(pca_components[:, 0], pca_components[:, 1], c=self.y,
↳cmap='viridis')
    plt.colorbar(label='Median Value')
    plt.xlabel('Principal Component 1')
    plt.ylabel('Principal Component 2')
    plt.title('PCA Visualization (2 components)')
    plt.show()

def summarize_data(self):
    """

```



```

        Summarize data
        """
        print(self.data.describe())

    def show_r2(self, y_test, y_pred):
        """
        show the r2 plot
        """
        # calc r2 score
        r2 = r2_score(y_test, y_pred)
        print(f"R2 Score: {r2:.4f}")

        # Plot y_test vs y_pred
        plt.figure(figsize=(8, 6))
        plt.scatter(y_test, y_pred, edgecolor='k', alpha=0.7, label='Predicted_
↪Vals vs Actual Vals')
        plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
↪'r--', lw=2, label='Line of Perfect fit')
        plt.xlabel('Actual Values')
        plt.ylabel('Predicted Values')
        plt.title(f'R2 Correlation: {r2:.4f}')
        plt.legend()
        plt.grid(True)
        plt.show()

```

```

[319]: # Load the dataset
df = pd.read_csv("HousingData.csv")

```

```

[320]: # Initialize the EDA class
eda = BostonHousingEDA(df)

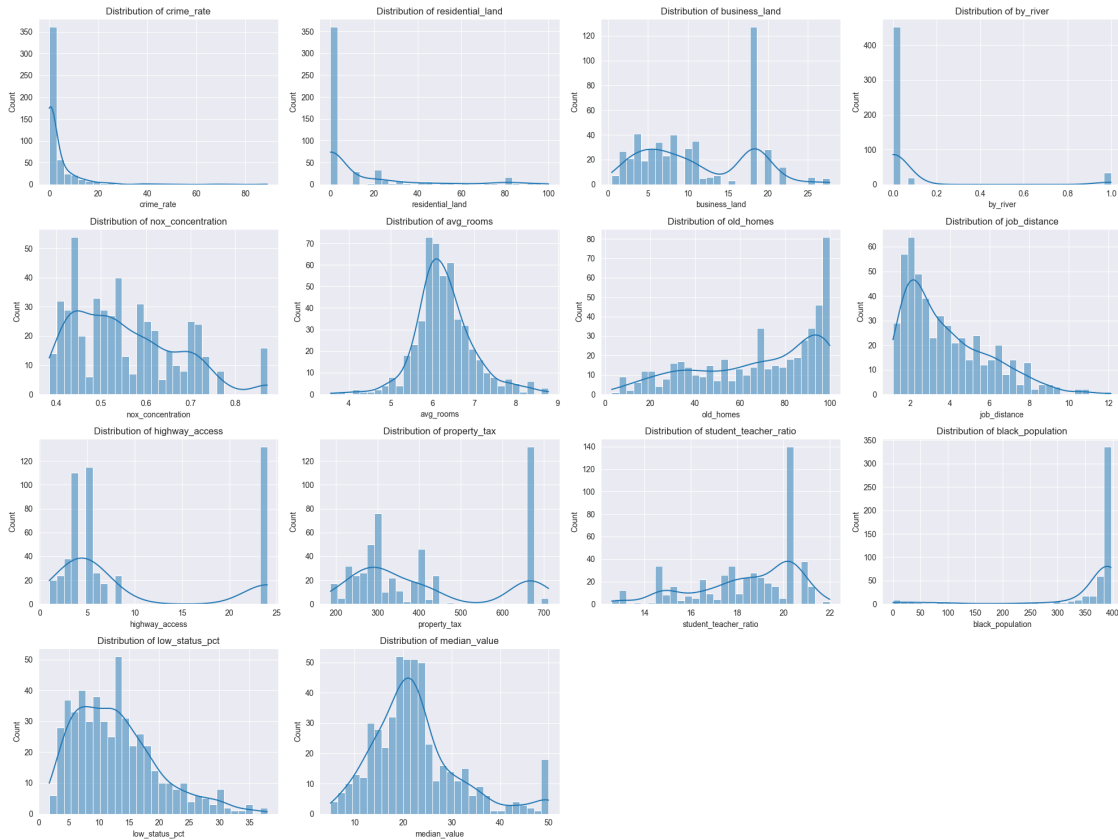
# Handle missing values
eda.handle_missing_values(strategy='mean')

```


```

[321]: # Visualize the data distribution
eda.visualize_data_distribution()

```



```
[322]: # Visualize the correlation matrix
eda.visualize_correlation_matrix()
```



### 1.1 3. Data Preprocessing:

```
[323]: eda.summarize_data()
```

	crime_rate	residential_land	business_land	by_river \
count	506.000000	506.000000	506.000000	506.000000
mean	3.611874	11.211934	11.083992	0.069959
std	8.545770	22.921051	6.699165	0.250233
min	0.006320	0.000000	0.460000	0.000000
25%	0.083235	0.000000	5.190000	0.000000
50%	0.290250	0.000000	9.900000	0.000000
75%	3.611874	11.211934	18.100000	0.000000
max	88.976200	100.000000	27.740000	1.000000

	nox_concentration	avg_rooms	old_homes	job_distance \
count	506.000000	506.000000	506.000000	506.000000
mean	0.554695	6.284634	68.518519	3.795043
std	0.115878	0.702617	27.439466	2.105710
min	0.385000	3.561000	2.900000	1.129600
25%	0.449000	5.885500	45.925000	2.100175

50%	0.538000	6.208500	74.450000	3.207450
75%	0.624000	6.623500	93.575000	5.188425
max	0.871000	8.780000	100.000000	12.126500

	highway_access	property_tax	student_teacher_ratio	black_population \
count	506.000000	506.000000	506.000000	506.000000
mean	9.549407	408.237154	18.455534	356.674032
std	8.707259	168.537116	2.164946	91.294864
min	1.000000	187.000000	12.600000	0.320000
25%	4.000000	279.000000	17.400000	375.377500
50%	5.000000	330.000000	19.050000	391.440000
75%	24.000000	666.000000	20.200000	396.225000
max	24.000000	711.000000	22.000000	396.900000

	low_status_pct	median_value
count	506.000000	506.000000
mean	12.715432	22.532806
std	7.012739	9.197104
min	1.730000	5.000000
25%	7.230000	17.025000
50%	11.995000	21.200000
75%	16.570000	25.000000
max	37.970000	50.000000

```
[324]: # Handle multicollinearity
X, y = eda.handle_multicollinearity(threshold=0.80)

X.head()
```

The columns dropped due to MCL: ['property\_tax']  
 Remaining columns: ['crime\_rate', 'residential\_land', 'business\_land', 'by\_river', 'nox\_concentration', 'avg\_rooms', 'old\_homes', 'job\_distance', 'highway\_access', 'student\_teacher\_ratio', 'black\_population', 'low\_status\_pct']

```
[324]:
```

	crime_rate	residential_land	business_land	by_river	nox_concentration \
0	0.00632	18.0	2.31	0.0	0.538
1	0.02731	0.0	7.07	0.0	0.469
2	0.02729	0.0	7.07	0.0	0.469
3	0.03237	0.0	2.18	0.0	0.458
4	0.06905	0.0	2.18	0.0	0.458

	avg_rooms	old_homes	job_distance	highway_access	student_teacher_ratio \
0	6.575	65.2	4.0900	1	15.3
1	6.421	78.9	4.9671	2	17.8
2	7.185	61.1	4.9671	2	17.8
3	6.998	45.8	6.0622	3	18.7
4	7.147	54.2	6.0622	3	18.7

	black_population	low_status_pct
0	396.90	4.980000
1	396.90	9.140000
2	392.83	4.030000
3	394.63	2.940000
4	396.90	12.715432

[324]:

```
[325]: # scale data
X, y = eda.scale_data(method='minmax')

print(X.shape)

X.head()
```

(506, 12)

```
[325]: crime_rate residential_land business_land by_river nox_concentration \
0 0.000000 0.18 0.067815 0.0 0.314815
1 0.000236 0.00 0.242302 0.0 0.172840
2 0.000236 0.00 0.242302 0.0 0.172840
3 0.000293 0.00 0.063050 0.0 0.150206
4 0.000705 0.00 0.063050 0.0 0.150206

avg_rooms old_homes job_distance highway_access student_teacher_ratio \
0 0.577505 0.641607 0.269203 0.000000 0.287234
1 0.547998 0.782698 0.348962 0.043478 0.553191
2 0.694386 0.599382 0.348962 0.043478 0.553191
3 0.658555 0.441813 0.448545 0.086957 0.648936
4 0.687105 0.528321 0.448545 0.086957 0.648936

black_population low_status_pct
0 1.000000 0.089680
1 1.000000 0.204470
2 0.989737 0.063466
3 0.994276 0.033389
4 1.000000 0.303130
```

[326]: print(X.head())

	crime_rate	residential_land	business_land	by_river	nox_concentration	\
0	0.000000	0.18	0.067815	0.0	0.314815	
1	0.000236	0.00	0.242302	0.0	0.172840	
2	0.000236	0.00	0.242302	0.0	0.172840	
3	0.000293	0.00	0.063050	0.0	0.150206	
4	0.000705	0.00	0.063050	0.0	0.150206	

	avg_rooms	old_homes	job_distance	highway_access	student_teacher_ratio	\
0	0.577505	0.641607	0.269203	0.000000	0.287234	
1	0.547998	0.782698	0.348962	0.043478	0.553191	
2	0.694386	0.599382	0.348962	0.043478	0.553191	
3	0.658555	0.441813	0.448545	0.086957	0.648936	
4	0.687105	0.528321	0.448545	0.086957	0.648936	


  

	black_population	low_status_pct
0	1.000000	0.089680
1	1.000000	0.204470
2	0.989737	0.063466
3	0.994276	0.033389
4	1.000000	0.303130

```
[327]: len(y)
```

```
[327]: 506
```

```
[328]: # visualize again
eda.pca_visualization(n_components=2)
```



```
[329]: # describe
eda.summarize_data()
```

	crime_rate	residential_land	business_land	by_river	\
count	506.000000	506.000000	506.000000	506.000000	
mean	3.611874	11.211934	11.083992	0.069959	
std	8.545770	22.921051	6.699165	0.250233	
min	0.006320	0.000000	0.460000	0.000000	
25%	0.083235	0.000000	5.190000	0.000000	
50%	0.290250	0.000000	9.900000	0.000000	
75%	3.611874	11.211934	18.100000	0.000000	
max	88.976200	100.000000	27.740000	1.000000	

	nox_concentration	avg_rooms	old_homes	job_distance	\
count	506.000000	506.000000	506.000000	506.000000	
mean	0.554695	6.284634	68.518519	3.795043	
std	0.115878	0.702617	27.439466	2.105710	
min	0.385000	3.561000	2.900000	1.129600	
25%	0.449000	5.885500	45.925000	2.100175	
50%	0.538000	6.208500	74.450000	3.207450	
75%	0.624000	6.623500	93.575000	5.188425	
max	0.871000	8.780000	100.000000	12.126500	

	highway_access	property_tax	student_teacher_ratio	black_population	\
count	506.000000	506.000000	506.000000	506.000000	
mean	9.549407	408.237154	18.455534	356.674032	
std	8.707259	168.537116	2.164946	91.294864	
min	1.000000	187.000000	12.600000	0.320000	
25%	4.000000	279.000000	17.400000	375.377500	
50%	5.000000	330.000000	19.050000	391.440000	
75%	24.000000	666.000000	20.200000	396.225000	
max	24.000000	711.000000	22.000000	396.900000	

	low_status_pct	median_value
count	506.000000	506.000000
mean	12.715432	22.532806
std	7.012739	9.197104
min	1.730000	5.000000
25%	7.230000	17.025000
50%	11.995000	21.200000
75%	16.570000	25.000000
max	37.970000	50.000000

```
[330]: from sklearn.model_selection import train_test_split

# Split the X and y from each other
```

```
# Test this on different scalers to see what works best
X, y = eda.scale_data(method='minmax')
```

```
[331]: # Split the data in training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[332]: print("X_train.shape", X_train.shape)
print("X_test.shape", X_test.shape)
print("y_train.shape", y_train.shape)
print("y_test.shape", y_test.shape)
```

```
X_train.shape (404, 12)
X_test.shape (102, 12)
y_train.shape (404,)
y_test.shape (102,)
```

```
[333]: X_train.head()
```

```
[333]:      crime_rate  residential_land  business_land  by_river  nox_concentration  \
477    0.168788         0.000000         0.646628         0.0         0.471193
15     0.006981         0.000000         0.281525         0.0         0.314815
332    0.000319         0.112119         0.205279         0.0         0.108848
423    0.079174         0.000000         0.646628         0.0         0.471193
19     0.008087         0.000000         0.281525         0.0         0.314815

      avg_rooms  old_homes  job_distance  highway_access  \
477    0.333972    0.972194     0.088307         1.000000
15     0.435524    0.552008     0.306359         0.130435
332    0.473271    0.210093     0.501150         0.000000
423    0.487066    0.675783     0.081132         1.000000
19     0.415022    0.685891     0.242514         0.130435

      student_teacher_ratio  black_population  low_status_pct
477             0.808511             0.880428         0.639625
15              0.893617             0.996772         0.185982
332              0.457447             0.912628         0.168322
423              0.808511             0.005547         0.594923
19              0.893617             0.984997         0.263521
```

```
[334]: y_train
```

```
[334]: 477    12.0
15     19.9
332    19.4
423    13.4
19     18.2
```



```

...
106    19.5
270    21.1
348    24.5
435    13.4
102    18.6
Name: median_value, Length: 404, dtype: float64

```

#### 1.1.1 4. OLS regression:

```

[335]: import numpy as np

class OLSRegression:
    def __init__(self):
        self.theta = None

    def fit(self, X, y):
        """
        Fit method that computes OLS
        """
        # add bias term to X
        ones_column = np.ones((X.shape[0], 1))
        X_bias = np.concatenate((ones_column, X), axis=1)

        # Compute  $X^T$ 
        X_transpose = X_bias.T

        # multiply  $X^T$  by X
        X_transpose_X = X_transpose.dot(X_bias)

        # inverse of  $(X^T * X)$ 
        X_transpose_X_inv = np.linalg.inv(X_transpose_X)

        # multiply the inverse by  $X^T$ 
        X_transpose_X_inv_X_transpose = X_transpose_X_inv.dot(X_transpose)

        # multiply by y to get the coefficients
        self.theta = X_transpose_X_inv_X_transpose.dot(y)

    def predict(self, X):
        """
        Make predictions
        """
        # add bias term to X
        ones_column = np.ones((X.shape[0], 1))
        X_bias = np.concatenate((ones_column, X), axis=1)

```

```
# multiply X_bias by theta
y_pred = X_bias.dot(self.theta)

# Return the pred values
return y_pred
```

```
[336]: from sklearn.metrics import mean_squared_error, r2_score
```

```
# Initialize and train the OLS regression model
ols_model = OLSRegression()

# Fit the model on training data
ols_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = ols_model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)


print(f"OLS Regression Mean Squared Error: {mse}")
print(f"OLS Regression R2 Score: {r2}")

eda.show_r2(y_test, y_pred)
```

OLS Regression Mean Squared Error: 25.812604382189356

OLS Regression R2 Score: 0.6480120993382052

R2 Score: 0.6480



```
[337]: from sklearn.linear_model import LinearRegression

# Testing Linear Regression from SKLearn to compare

# Init the linreg model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)


# predict
y_pred_lr = lr_model.predict(X_test)

# Evaluate like previously using MSE and R2
mse_lr = mean_squared_error(y_test, y_pred_lr)
r2_lr = r2_score(y_test, y_pred_lr)

print(f"SLK LinearRegression Mean Squared Error: {mse_lr}")
print(f"SKL LinearRegression R2 Score: {r2_lr}")

eda.show_r2(y_test, y_pred)
```

SLK LinearRegression Mean Squared Error: 25.812604382189612  
SKL LinearRegression R2 Score: 0.6480120993382017  
R2 Score: 0.6480



### 1.1.2 5. Ridge Regression:

```
[338]: class RidgeRegression:
    def __init__(self, alpha=1.0):
        """
        Init the Ridge Regression model
        """
        self.alpha = alpha
        self.theta = None

    def fit(self, X, y):
        """
        Fit method for Ridge Regression using the normal equation
        """
        # add bias term to X
```

```

X_bias = np.c_[np.ones((X.shape[0], 1)), X]

# Calculate theta using the normal equation with L2 reg
identity = np.eye(X_bias.shape[1])
identity[0, 0] = 0
self.theta = np.linalg.inv(X_bias.T.dot(X_bias) + self.alpha * ↳
identity).dot(X_bias.T).dot(y)

def predict(self, X):
    """
    predict method for ridge reg model
    """
    # add bias term (intercept)
    X_bias = np.c_[np.ones((X.shape[0], 1)), X]
    return X_bias.dot(self.theta)

```

```

[339]: alphas = [0.5, 1, 1.5, 2]

#iterate over multiple alpha vals
for alpha in alphas:
    print("-----")
    print(f"\nAlpha = {alpha}:")
    ridge = RidgeRegression(alpha=alpha)
    ridge.fit(X_train, y_train)
    y_pred_ridge = ridge.predict(X_test)
    print(f"Ridge MSE: {mean_squared_error(y_test, y_pred_ridge)}")
    print(f"Ridge R2: {r2_score(y_test, y_pred_ridge)}")

```

```

-----

Alpha = 0.5:
Ridge MSE: 25.19079436852606
Ridge R2: 0.6564912747859506
-----

```

```

Alpha = 1:
Ridge MSE: 24.832196260159254
Ridge R2: 0.6613812189960132
-----

```

```

Alpha = 1.5:
Ridge MSE: 24.63624181400551
Ridge R2: 0.6640533086893179
-----

```

```

Alpha = 2:
Ridge MSE: 24.548498714424547
Ridge R2: 0.665249798162514

```

### 1.1.3 5. Lasso Regression:

```
[340]: class LassoRegression:
    def __init__(self, alpha=1.0, n_iterations=1000, learning_rate=0.01):
        """
        Init the Lasso Regression model
        """
        self.alpha = alpha
        self.n_iterations = n_iterations
        self.learning_rate = learning_rate
        self.theta = None

    def fit(self, X, y):
        """
        Fir method using coordinate descent
        """
        # add bias term to X
        X_bias = np.c_[np.ones((X.shape[0], 1)), X]

        # set samples and feats
        n_samples, n_features = X_bias.shape

        # init theta
        self.theta = np.zeros(n_features)

        # implement gradient descent
        for _ in range(self.n_iterations):

            # set pred
            y_pred = X_bias.dot(self.theta)

            # set gradient via bias times dot
            gradients = X_bias.T.dot(y_pred - y) / n_samples

            # update theta with L1 reg
            for j in range(1, n_features):
                self.theta[j] -= self.learning_rate * (gradients[j] + self.
↪alpha * np.sign(self.theta[j]))

            # update intercept separately
            self.theta[0] -= self.learning_rate * gradients[0]

    def predict(self, X):
        """
        pred new data with Lasso reg
        """
        # Add bias term
```

```

X_bias = np.c_[np.ones((X.shape[0], 1)), X]

# return product
return X_bias.dot(self.theta)

```

```

[341]: from sklearn.metrics import mean_squared_error, r2_score

#iterate over multiple alpha vals
alphas = [0.5, 1, 1.5, 2]
for alpha in alphas:
    print("-----")
    print(f"\nalpha = {alpha}:")
    lasso = LassoRegression(alpha=alpha, learning_rate=0.01, n_iterations=1000)
    lasso.fit(X_train, y_train)
    y_pred_lasso = lasso.predict(X_test)
    print(f"Lasso MSE: {mean_squared_error(y_test, y_pred_lasso)}")
    print(f"Lasso R2: {r2_score(y_test, y_pred_lasso)}")

```

```

-----

alpha = 0.5:
Lasso MSE: 49.00844114265968
Lasso R2: 0.33170717463850197
-----

```

```

alpha = 1:
Lasso MSE: 62.05883125966329
Lasso R2: 0.15374840100653564
-----

```

```

alpha = 1.5:
Lasso MSE: 66.47302374946402
Lasso R2: 0.09355523627984508
-----

```

```

alpha = 2:
Lasso MSE: 70.73738742806061
Lasso R2: 0.03540517917350694

```

```

[342]: from sklearn.linear_model import Ridge, Lasso

alphas = [0.5, 0.75, 1, 1.5, 2]

# iterate over multiple alpha vals
for alpha in alphas:
    # ridge model
    print(f"\nSklearn Ridge alpha = {alpha}:")
    sklearn_ridge = Ridge(alpha=alpha)

```

```

sklearn_ridge.fit(X_train, y_train)
y_pred_ridge_sklearn = sklearn_ridge.predict(X_test)
print(f"Ridge MSE: {mean_squared_error(y_test, y_pred_ridge_sklearn)}")
print(f"Ridge R2: {r2_score(y_test, y_pred_ridge_sklearn)}")

# lasso model
print(f"\nSklearn Lasso alpha = {alpha}:")
sklearn_lasso = Lasso(alpha=alpha, max_iter=1000)
sklearn_lasso.fit(X_train, y_train)
y_pred_lasso_sklearn = sklearn_lasso.predict(X_test)
print(f"Lasso MSE: {mean_squared_error(y_test, y_pred_lasso_sklearn)}")
print(f"Lasso R2: {r2_score(y_test, y_pred_lasso_sklearn)}")

```

Sklearn Ridge alpha = 0.5:  
Ridge MSE: 25.19079436852541  
Ridge R<sup>2</sup>: 0.6564912747859595

Sklearn Lasso alpha = 0.5:  
Lasso MSE: 33.62830965344012  
Lasso R<sup>2</sup>: 0.5414349539294651

Sklearn Ridge alpha = 0.75:  
Ridge MSE: 24.986408159234436  
Ridge R<sup>2</sup>: 0.6592783423622244

Sklearn Lasso alpha = 0.75:  
Lasso MSE: 44.60185159749377  
Lasso R<sup>2</sup>: 0.39179666348339326

Sklearn Ridge alpha = 1:  
Ridge MSE: 24.83219626015898  
Ridge R<sup>2</sup>: 0.6613812189960169

Sklearn Lasso alpha = 1:  
Lasso MSE: 57.565892476652685  
Lasso R<sup>2</sup>: 0.21501537223571798

Sklearn Ridge alpha = 1.5:  
Ridge MSE: 24.636241814005345  
Ridge R<sup>2</sup>: 0.6640533086893202

Sklearn Lasso alpha = 1.5:  
Lasso MSE: 75.04543037399255  
Lasso R<sup>2</sup>: -0.023340500652033302

Sklearn Ridge alpha = 2:  
Ridge MSE: 24.548498714424312



Ridge  $R^2$ : 0.6652497981625172

Sklearn Lasso alpha = 2:

Lasso MSE: 75.04543037399255

Lasso  $R^2$ : -0.023340500652033302

## 1.2 6. Compare Results:

```
[343]: import pandas as pd
from sklearn.preprocessing import StandardScaler, MinMaxScaler, PowerTransformer
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.linear_model import Ridge, Lasso

# Initialize and train the OLS regression model
def evaluate_models(X_train, X_test, y_train, y_test, alpha=None):
    # Initialize result dictionary
    results = {
        "Model": [],
        "alpha": [],
        "scaler": [],
        "Custom R2": [],
        "Custom MSE": [],
        "Sklearn R2": [],
        "Sklearn MSE": []
    }

    # custom OLS
    if alpha is None:
        ols_model = OLSRegression()
        ols_model.fit(X_train, y_train)
        y_pred_ols = ols_model.predict(X_test)
        mse_ols = mean_squared_error(y_test, y_pred_ols)
        r2_ols = r2_score(y_test, y_pred_ols)

        # Store OLS
        results["Model"].append("OLS")
        results["alpha"].append(None)
        results["scaler"].append(current_scaler_name)
        results["Custom R2"].append(r2_ols)
        results["Custom MSE"].append(mse_ols)
        results["Sklearn R2"].append(None)
        results["Sklearn MSE"].append(None)

    else:
        #Custom ridge
        custom_ridge = RidgeRegression(alpha=alpha)
        custom_ridge.fit(X_train, y_train)
```

```

y_pred_custom_ridge = custom_ridge.predict(X_test)
mse_custom_ridge = mean_squared_error(y_test, y_pred_custom_ridge)
r2_custom_ridge = r2_score(y_test, y_pred_custom_ridge)

# SKL Ridge
sklearn_ridge = Ridge(alpha=alpha)
sklearn_ridge.fit(X_train, y_train)
y_pred_sklearn_ridge = sklearn_ridge.predict(X_test)
mse_sklearn_ridge = mean_squared_error(y_test, y_pred_sklearn_ridge)
r2_sklearn_ridge = r2_score(y_test, y_pred_sklearn_ridge)

# Save results
results["Model"].append("Ridge")
results["alpha"].append(alpha)
results["scaler"].append(current_scaler_name)
results["Custom R2"].append(r2_custom_ridge)
results["Custom MSE"].append(mse_custom_ridge)
results["Sklearn R2"].append(r2_sklearn_ridge)
results["Sklearn MSE"].append(mse_sklearn_ridge)

# custom lasso
custom_lasso = LassoRegression(alpha=alpha, learning_rate=0.01,
↪n_iterations=1000)
custom_lasso.fit(X_train, y_train)
y_pred_custom_lasso = custom_lasso.predict(X_test)
mse_custom_lasso = mean_squared_error(y_test, y_pred_custom_lasso)
r2_custom_lasso = r2_score(y_test, y_pred_custom_lasso)

# sklearn lasso
sklearn_lasso = Lasso(alpha=alpha, max_iter=1000)
sklearn_lasso.fit(X_train, y_train)
y_pred_sklearn_lasso = sklearn_lasso.predict(X_test)
mse_sklearn_lasso = mean_squared_error(y_test, y_pred_sklearn_lasso)
r2_sklearn_lasso = r2_score(y_test, y_pred_sklearn_lasso)

# store results
results["Model"].append("Lasso")
results["alpha"].append(alpha)
results["scaler"].append(current_scaler_name)
results["Custom R2"].append(r2_custom_lasso)
results["Custom MSE"].append(mse_custom_lasso)
results["Sklearn R2"].append(r2_sklearn_lasso)
results["Sklearn MSE"].append(mse_sklearn_lasso)

return results

scalers = {

```

```

    "StandardScaler": StandardScaler(),
    "MinMaxScaler": MinMaxScaler(),
    "PowerTransformer": PowerTransformer()
}

# alpha values
alphas = [0.5, 1, 1.5, 2]

# Create DF
comparison_df = pd.DataFrame()

# Scaler mapping for original class
scaler_mapping = {
    "StandardScaler": "standard",
    "MinMaxScaler": "minmax",
    "PowerTransformer": "power"
}

# Loop through each scaler
for current_scaler_name, current_scaler in scalers.items():
    # Scale data
    X, y = eda.scale_data(method=scaler_mapping[current_scaler_name])

    # split to train and test
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
        random_state=42)

    # evaluate OLS
    ols_results = evaluate_models(X_train, X_test, y_train, y_test, alpha=None)
    comparison_df = pd.concat([comparison_df, pd.DataFrame(ols_results)],
        ignore_index=True)

    # loop through alphas and calc ridge and lasso
    for alpha in alphas:
        current_results = evaluate_models(X_train, X_test, y_train, y_test,
            alpha)
        comparison_df = pd.concat([comparison_df, pd.
            DataFrame(current_results)], ignore_index=True)

# Show results
comparison_df

```

/var/folders/r4/9ncp61z543v94\_5sltg11\_bm0000gn/T/ipykernel\_24506/2601374963.py:19: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```

comparison_df = pd.concat([comparison_df, pd.DataFrame(current_results)],
ignore_index=True)
/var/folders/r4/9ncp61z543v94_5sltgt11_bm0000gn/T/ipykernel_24506/2601374963.py:1
14: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA
entries is deprecated. In a future version, this will no longer exclude empty or
all-NA columns when determining the result dtypes. To retain the old behavior,
exclude the relevant entries before the concat operation.
comparison_df = pd.concat([comparison_df, pd.DataFrame(ols_results)],
ignore_index=True)
/var/folders/r4/9ncp61z543v94_5sltgt11_bm0000gn/T/ipykernel_24506/2601374963.py:1
14: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA
entries is deprecated. In a future version, this will no longer exclude empty or
all-NA columns when determining the result dtypes. To retain the old behavior,
exclude the relevant entries before the concat operation.
comparison_df = pd.concat([comparison_df, pd.DataFrame(ols_results)],
ignore_index=True)

```

```

[343]:

```

	Model	alpha	scaler	Custom R2	Custom MSE	Sklearn R2	\
0	OLS	NaN	StandardScaler	0.648012	25.812604	NaN	
1	Ridge	0.5	StandardScaler	0.648107	25.805644	0.648107	
2	Lasso	0.5	StandardScaler	0.621555	27.752815	0.621365	
3	Ridge	1.0	StandardScaler	0.648199	25.798907	0.648199	
4	Lasso	1.0	StandardScaler	0.619499	27.903557	0.619385	
5	Ridge	1.5	StandardScaler	0.648288	25.792387	0.648288	
6	Lasso	1.5	StandardScaler	0.606851	28.831112	0.606378	
7	Ridge	2.0	StandardScaler	0.648374	25.786075	0.648374	
8	Lasso	2.0	StandardScaler	0.580056	30.796118	0.580000	
9	OLS	NaN	MinMaxScaler	0.648012	25.812604	NaN	
10	Ridge	0.5	MinMaxScaler	0.656491	25.190794	0.656491	
11	Lasso	0.5	MinMaxScaler	0.331707	49.008441	0.541435	
12	Ridge	1.0	MinMaxScaler	0.661381	24.832196	0.661381	
13	Lasso	1.0	MinMaxScaler	0.153748	62.058831	0.215015	
14	Ridge	1.5	MinMaxScaler	0.664053	24.636242	0.664053	
15	Lasso	1.5	MinMaxScaler	0.093555	66.473024	-0.023341	
16	Ridge	2.0	MinMaxScaler	0.665250	24.548499	0.665250	
17	Lasso	2.0	MinMaxScaler	0.035405	70.737387	-0.023341	
18	OLS	NaN	PowerTransformer	0.702784	21.795955	NaN	
19	Ridge	0.5	PowerTransformer	0.702713	21.801170	0.702713	
20	Lasso	0.5	PowerTransformer	0.678410	23.583409	0.680138	
21	Ridge	1.0	PowerTransformer	0.702638	21.806670	0.702638	
22	Lasso	1.0	PowerTransformer	0.668776	24.289904	0.669681	
23	Ridge	1.5	PowerTransformer	0.702560	21.812436	0.702560	
24	Lasso	1.5	PowerTransformer	0.655367	25.273234	0.655860	
25	Ridge	2.0	PowerTransformer	0.702478	21.818449	0.702478	
26	Lasso	2.0	PowerTransformer	0.630445	27.100889	0.631805	

Sklearn MSE

0	NaN
1	25.805644
2	27.766770
3	25.798907
4	27.911920
5	25.792387
6	28.865816
7	25.786075
8	30.800179
9	NaN
10	25.190794
11	33.628310
12	24.832196
13	57.565892
14	24.636242
15	75.045430
16	24.548499
17	75.045430
18	NaN
19	21.801170
20	23.456674
21	21.806670
22	24.223564
23	21.812436
24	25.237088
25	21.818449
26	27.001166

- The custom models and sklearn models show nearly identical robustness and performance in terms of correlation R2 and MSE confirming the correctness of the custom implementations above and correctness of code.
- PowerTransformer scaler consistently shows the best performance across both Ridge and Lasso. We observe this with the highest R2 correlations and lowest MSE values, which means PowerTransformer is doing a better job at transforming the data making the models more effective
- StandardScaler and MinMaxScaler underperform slightly compared to PowerTransformer, showing lower R2 and higher MSE values across the models.
- While both StandardScaler and MinMaxScaler can stabilize the data they don't seem to provide the optimal transformation that PowerTransformer does in this dataset as seen above in the results
- Ridge regression is more stable across different values of alpha. As alpha increases there's only a slight decrease in R2 and a slightly increase in MSE in the different iterations in the loop which shows Ridge is handling regularization well due to its L2 penalty. This consistency makes Ridge reliable when working with different levels of regularization.
- Lasso regression is more sensitive to changes in alpha, showing significant performance degra-

dation (we see lower R2, higher MSE) as alpha increases. This sensitivity comes from the L1 regularization which aggressively shrinks coefficients and can set them to zero, leading to underfitting if alpha becomes too large like we saw in class

- Ridge generally outperforms Lasso in terms of R2 and MSE, particularly when using PowerTransformer. Ridge retains more information by shrinking coefficients without eliminating them, which is especially useful when most features contribute to the predictive power of the model.
- Lasso is more aggressive in shrinking coefficients, which can result in underfitting as alpha increases. This is especially true when some features are important but are driven to zero because of the model. This explains why Lasso struggles as alpha goes up leading to worse performance compared to Ridge like we see in the table above.
- Regularization has a greater impact on Lasso due to L1 red, while Ridge L2 regularization offers a more balanced trade-off between bias and variance. This is particularly true in datasets with many relevant features like this one, where L1's feature selection might lead to underfitting.
- OLS regression performance is very close to Ridge with low regularization, especially when using PowerTransformer. OLS shows strong R2 and MSE values showing that without regularization, Ridge tends to behave similarly to OLS in terms of model complexity and performance.
- MinMaxScaler shows the worst performance with Lasso regression, particularly at higher alpha values, with significantly bad results in R2 and MSE. This may indicate that MinMaxScaler doesn't transform the data effectively for Lasso, leading to underfitting at even moderate levels of regularization
- PowerTransformer proves to be the most effective scaler for both Ridge and Lasso, providing stability and better model performance across various alpha values, especially when you compare to StandardScaler and MinMaxScaler.