

assignment3_task2

October 31, 2024

1 Assignment 3 - Task 2

1.0.1 Explore the solution space for circles, and moons. Datasets also provided in code of Module 6 (25)

1. Find best kernel function
2. Visualize and explain solution

Contents:

1. Import Libraries
2. Create Explorer Class
3. Creating Circles Dataset and Explore
4. Creating Moons Dataset and Explore
5. Evaluate the Kernels and Noise
6. Explain Solution

1.0.2 1. Import Libraries

```
[44]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import make_moons, make_circles
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, roc_auc_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
```

1.0.3 2. Create Explorer Class

```
[55]: class KernelSVMExplorer:
    def __init__(self, kernels=['linear', 'poly', 'rbf', 'sigmoid']):
        """
        Initiallize the class with kernel options
        """
        self.kernels = kernels
```

```

def generate_dataset(self, dataset_name='moons', **kwargs):
    """
    generate dataset using the specified name to make it modular.
    """
    # pick dataset
    if dataset_name == 'moons':
        X, y = make_moons(**kwargs)
    elif dataset_name == 'circles':
        X, y = make_circles(**kwargs)
    else:
        raise ValueError("Error check input")
    return X, y

# 3. Visualize Dataset
def visualize_dataset(self, X, y, dataset_name):
    """
    plots the dataset with
    """
    # create figure
    plt.figure(figsize=(8, 6))
    plt.scatter(X[y == 0][:, 0], X[y == 0][:, 1], color='blue',
    ↪label='Class 0')
    plt.scatter(X[y == 1][:, 0], X[y == 1][:, 1], color='red', label='Class
    ↪1')

    #add labels and title
    plt.title(f"{dataset_name.capitalize()} Dataset")
    plt.xlabel("Feature 1")
    plt.ylabel("Feature 2")
    plt.legend()
    plt.grid(True)
    plt.show()

def evaluate_kernel(self, X_train, X_test, y_train, y_test, kernel):
    """
    train and evaluate SVM given the data and kernel
    """
    # create model and fit
    model = SVC(kernel=kernel, probability=True)
    model.fit(X_train, y_train)

    # get output
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)[:, 1] if kernel != 'linear' else
    ↪None

```

```

    # get metrics
    accuracy = accuracy_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_prob) if y_prob is not None else "N/A"
    return {"Kernel": kernel, "Accuracy": accuracy, "ROC-AUC": roc_auc}

def find_best_kernel(self, X, y):
    """
    Extends the evaluate_kernel method above to find the best kernel
    """

    # train test split
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪3, random_state=42)
    results = []

    # iterate over the kernels and evaluate each
    for kernel in self.kernels:
        metrics = self.evaluate_kernel(X_train, X_test, y_train, y_test,
↪kernel)

        # Handle "N/A" to avoid error
        if metrics["ROC-AUC"] == "N/A":
            metrics["ROC-AUC"] = 0.0

        results.append(metrics)

    # sort by acc, then ROC-AUC
    sorted_results = sorted(results, key=lambda x: (x['Accuracy'],
↪x['ROC-AUC']), reverse=True)
    return sorted_results[0]

def plot_decision_boundary(self, model, X, y, kernel, dataset_name):
    """
    Implementation from SKlearn to plot the decision boundary
    """

    # create plot
    plt.figure(figsize=(8, 6))
    plt.scatter(X[y == 0][:, 0], X[y == 0][:, 1], color='blue',
↪label='Class 0')
    plt.scatter(X[y == 1][:, 0], X[y == 1][:, 1], color='red', label='Class
↪1')

    # create mesh grid, code is based on SKLEARN's implementation
    x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
    y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5

```

```

        xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100), np.linspace(y_min,
↪y_max, 100))
        Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
        Z = Z.reshape(xx.shape)

        # Plot the title and legend
        plt.contourf(xx, yy, Z, alpha=0.2, cmap='coolwarm')
        plt.title(f"{dataset_name.capitalize()} - Kernel: {kernel}")
        plt.legend()
        plt.show()

def analyze_and_visualize(self, dataset_name='moons', **kwargs):
    """
    finds the best kernel and metrics automatically
    """
    # Split dataset
    X, y = self.generate_dataset(dataset_name, **kwargs)
    self.visualize_dataset(X, y, dataset_name)

    # Get best kernel
    best_kernel_metrics = self.find_best_kernel(X, y)
    print(f"Best Kernel for {dataset_name}:↵
↪{best_kernel_metrics['Kernel']}")
    print(f"Accuracy: {best_kernel_metrics['Accuracy']}, ROC-AUC:↵
↪{best_kernel_metrics['ROC-AUC']}")

    # Train final model
    best_model = SVC(kernel=best_kernel_metrics['Kernel'])
    best_model.fit(X, y)

    # Visualize decision boundary
    self.plot_decision_boundary(best_model, X, y,↵
↪best_kernel_metrics['Kernel'], dataset_name)

def evaluate_kernels_with_noise(self, dataset_names=['moons', 'circles'],↵
↪noise_levels=[0.1, 0.2, 0.3]):
    """
    evaluate each kernel and noise levels for both datasets. Return DF.
    """
    results = []

    # loop over the datasets
    for dataset_name in dataset_names:

        # loop over noises
        for noise in noise_levels:

```

```

        X, y = self.generate_dataset(dataset_name, noise=noise,
↪n_samples=200, random_state=1)
        X_train, X_test, y_train, y_test = train_test_split(X, y,
↪test_size=0.3, random_state=42)

        # loop over kernels
        for kernel in self.kernels:
            metrics = self.evaluate_kernel(X_train, X_test, y_train,
↪y_test, kernel)

            metrics.update({"Dataset": dataset_name, "Noise": noise})
            results.append(metrics)

        # create the DF and return
        df_results = pd.DataFrame(results)
        return df_results

    from sklearn.model_selection import cross_val_score

    def evaluate_kernels_with_noise_cross_validation(self,
↪dataset_names=['moons', 'circles'], noise_levels=[0.1, 0.2, 0.3], k=5):
        """
        eval kernels across noise levels but with f fold cross validation
        """
        results = []

        # loop datasets
        for dataset_name in dataset_names:

            # loop noise
            for noise in noise_levels:
                X, y = self.generate_dataset(dataset_name, noise=noise,
↪n_samples=200, random_state=1)

                # loop kernels
                for kernel in self.kernels:
                    model = SVC(kernel=kernel, probability=True, random_state=1)

                    # k-fold CV
                    scores = cross_val_score(model, X, y, cv=k,
↪scoring='accuracy')
                    mean_accuracy = scores.mean()

                    # capture metrics
                    metrics = {
                        "Kernel": kernel,
                        "Dataset": dataset_name,

```

```

        "Noise": noise,
        "k-Fold CV Accuracy": mean_accuracy
    }
    results.append(metrics)

df_results = pd.DataFrame(results)
return df_results

```

1.0.4 3. Creating Circles Dataset and Explore

```


[56]: # Instantiate the class
explorer = KernelSVMExplorer()

```


```

[57]: # Generate and analyze the circles dataset
explorer.analyze_and_visualize('circles', n_samples=200, factor=0.5, noise=0.1,
    ↪ random_state=1)


```



Best Kernel for circles: rbf
 Accuracy: 0.9833333333333333, ROC-AUC: 1.0




```
[58]: explorer.analyze_and_visualize('circles', n_samples=200, factor=0.5, noise=0.2, random_state=1)
```




Best Kernel for circles: rbf


Accuracy: 0.8833333333333333, ROC-AUC: 0.9927048260381593



```
[59]: explorer.analyze_and_visualize('circles', n_samples=200, factor=0.5, noise=0.3, random_state=1)
```




Best Kernel for circles: rbf
Accuracy: 0.75, ROC-AUC: 0.9393939393939394

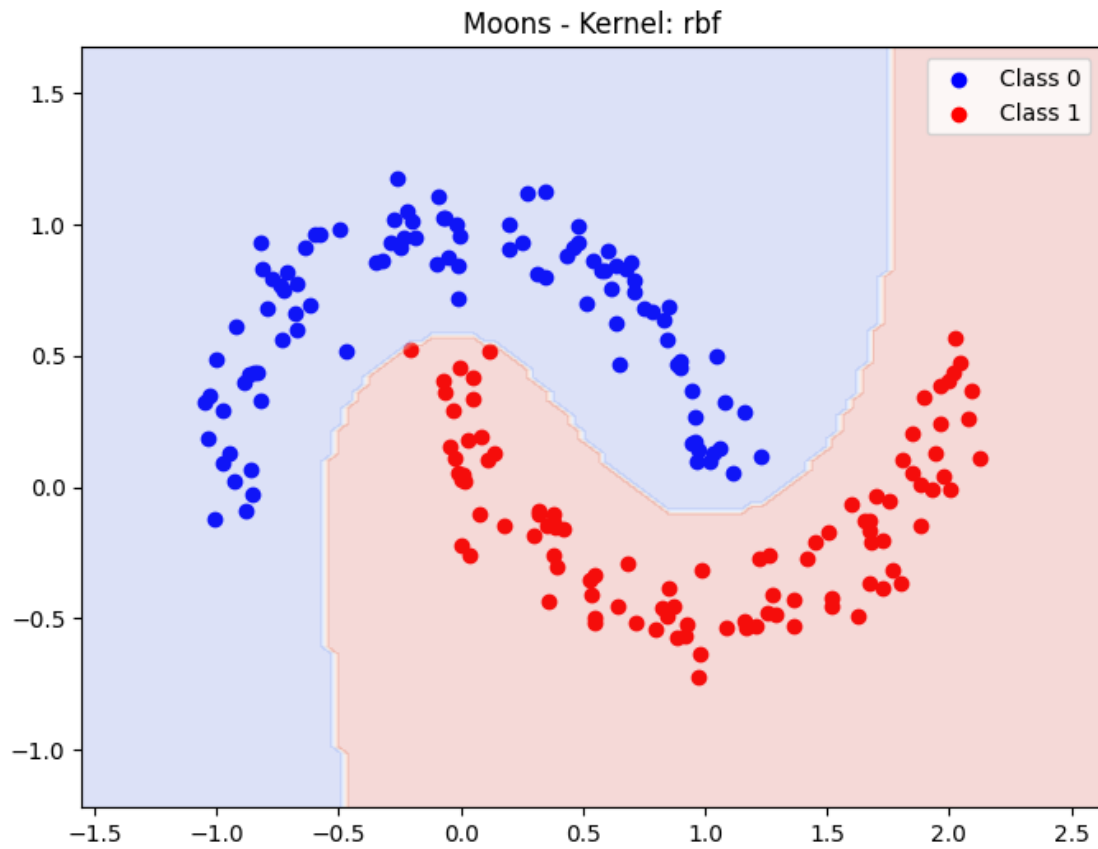


1.0.5 4. Creating Moons Dataset and Explore


```
[60]: # generate and analyze moons dataset
explorer.analyze_and_visualize('moons', n_samples=200, noise=0.1,
    ↪ random_state=1)
```




Best Kernel for moons: rbf
Accuracy: 1.0, ROC-AUC: 1.0




```
[61]: explorer.analyze_and_visualize('moons', n_samples=200, noise=0.2,   
    ↪ random_state=1)
```



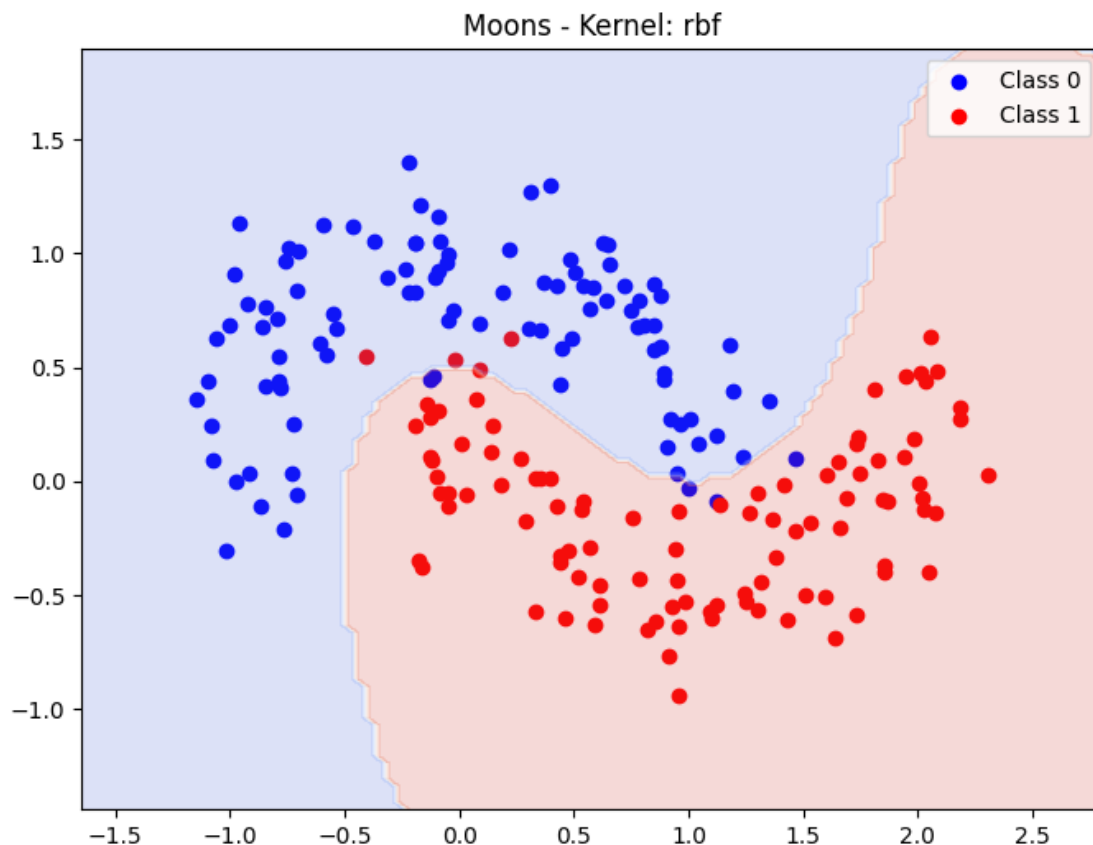
Best Kernel for moons: rbf
Accuracy: 0.95, ROC-AUC: 0.9932659932659933



```
[62]: explorer.analyze_and_visualize('moons', n_samples=200, noise=0.2,   
    ↪ random_state=1)
```



Best Kernel for moons: rbf
Accuracy: 0.95, ROC-AUC: 0.9932659932659933



1.0.6 5. Evaluate the Kernels and Noise

```
[63]: df_results = explorer.evaluate_kernels_with_noise()
      print(df_results)
```

	Kernel	Accuracy	ROC-AUC	Dataset	Noise
0	linear	0.883333	N/A	moons	0.1
1	poly	0.916667	0.974186	moons	0.1
2	rbf	1.000000	1.0	moons	0.1
3	sigmoid	0.600000	0.784512	moons	0.1
4	linear	0.866667	N/A	moons	0.2
5	poly	0.883333	0.962963	moons	0.2
6	rbf	0.950000	0.993266	moons	0.2
7	sigmoid	0.666667	0.808081	moons	0.2
8	linear	0.866667	N/A	moons	0.3
9	poly	0.883333	0.956229	moons	0.3
10	rbf	0.933333	0.969697	moons	0.3
11	sigmoid	0.666667	0.813692	moons	0.3
12	linear	0.450000	N/A	circles	0.1
13	poly	0.450000	0.514029	circles	0.1

14	rbf	0.800000	0.975309	circles	0.1
15	sigmoid	0.500000	0.483726	circles	0.1
16	linear	0.450000	N/A	circles	0.2
17	poly	0.500000	0.59147	circles	0.2
18	rbf	0.600000	0.804714	circles	0.2
19	sigmoid	0.533333	0.476992	circles	0.2
20	linear	0.450000	N/A	circles	0.3
21	poly	0.466667	0.613917	circles	0.3
22	rbf	0.533333	0.685746	circles	0.3
23	sigmoid	0.500000	0.508418	circles	0.3

```
[64]: # Apply with kfold cross validation
df_cross_val_results = explorer.evaluate_kernels_with_noise_cross_validation(
    dataset_names=['moons', 'circles'],
    noise_levels=[0.1, 0.2, 0.3],
    k=5
)

print("CV Results:\n", df_cross_val_results)
```

Cross-Validated Results:

	Kernel	Dataset	Noise	k-Fold CV Accuracy
0	linear	moons	0.1	0.880
1	poly	moons	0.1	0.925
2	rbf	moons	0.1	1.000
3	sigmoid	moons	0.1	0.655
4	linear	moons	0.2	0.875
5	poly	moons	0.2	0.880
6	rbf	moons	0.2	0.955
7	sigmoid	moons	0.2	0.630
8	linear	moons	0.3	0.850
9	poly	moons	0.3	0.835
10	rbf	moons	0.3	0.910
11	sigmoid	moons	0.3	0.630
12	linear	circles	0.1	0.385
13	poly	circles	0.1	0.470
14	rbf	circles	0.1	0.830
15	sigmoid	circles	0.1	0.530
16	linear	circles	0.2	0.400
17	poly	circles	0.2	0.585
18	rbf	circles	0.2	0.635
19	sigmoid	circles	0.2	0.430
20	linear	circles	0.3	0.415
21	poly	circles	0.3	0.575
22	rbf	circles	0.3	0.560
23	sigmoid	circles	0.3	0.410

1.0.7 6. Explain Solution

- The first observation is that the RBF kernel consistently outperformed other kernels. It achieved a pretty high accuracy and ROC-AUC, especially for the moons dataset. This infers that the RBF kernel is well-suited for handling non-linear boundaries present in both moons and circles.
- The moons dataset is generally easier to classify than the circles dataset, as shown by the higher accuracy and ROCAUC scores for most kernels and noise levels. Even with increased noise, the RBF kernel maintained pretty good performance for moons, suggesting it adapts well to moderately noisy data.
- In the circles dataset the RBF kernel performed significantly better than other kernels, but its performance declined with increased noise levels. We increased the noise gradually to see the differences there. This pattern shows that while the RBF kernel can capture circular boundaries, higher noise levels degrade its effectiveness.
- Linear and polynomial kernels struggled with both datasets as shown in the results above, particularly the circles dataset, where linear separability is not sufficient. The moons dataset, while slightly better used to these kernels, still required non-linear approaches for the best classification results. We used the noise here to really test the robustness of the models.