

assignment3_task3

October 31, 2024

1 Assignment 3 - Task 3

1.0.1 Train a simplest possible MLP to solve the same classification problems of task 2. (25)

1. Explain rationale behind your training process and design parameters
2. Compare SVM and MLP solution in terms of computational cost
3. Apply k-Fold cross validation

Contents:

1. Import Libraries
2. Generate Datasets
3. Train Simple MLP
4. Explain Rationale
5. Compare SVM and MLP
6. Apply k-Fold cross validation
7. Conclusion

1.0.2 1. Import Libraries:

```
[66]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import make_moons, make_circles
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.neural_network import MLPClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, roc_auc_score, recall_score, f1_score
import time

class ModelComparison:
    def __init__(self, hidden_layer_sizes=(5,), max_iter=10000,
        kernels=['linear', 'rbf']):
        """
```

```

Initialize with params for the mlp
"""

self.hidden_layer_sizes = hidden_layer_sizes
self.max_iter = max_iter
self.kernels = kernels

def generate_dataset(self, dataset_name='moons', **kwargs):
    """
    Generates dataset via name
    """
    if dataset_name == 'moons':
        X, y = make_moons(**kwargs)
    elif dataset_name == 'circles':
        X, y = make_circles(**kwargs)
    else:
        raise ValueError("error bad name")
    return X, y

def visualize_dataset(self, X, y, dataset_name):
    """
    plots dataset
    """
    plt.figure(figsize=(8, 6))
    plt.scatter(X[y == 0][:, 0], X[y == 0][:, 1], color='blue',
↪label='Class 0')
    plt.scatter(X[y == 1][:, 0], X[y == 1][:, 1], color='red', label='Class_
↪1')
    plt.title(f"{dataset_name.capitalize()} Dataset")
    plt.xlabel("Feature 1")
    plt.ylabel("Featere 2")
    plt.legend()
    plt.grid(True)
    plt.show()

def train_mlp(self, X, y):
    """
    train mlp on dataset and return metrics
    """
    # split up
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪3, random_state=1)
    mlp = MLPClassifier(hidden_layer_sizes=self.hidden_layer_sizes,
↪max_iter=self.max_iter, random_state=1)

    # training time
    start_time = time.time()
    mlp.fit(X_train, y_train)

```

```

train_time = time.time() - start_time

# eval model
y_pred = mlp.predict(X_test)
y_prob = mlp.predict_proba(X_test)[: , 1]
accuracy = accuracy_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_prob)

return {"Model": "MLP", "Accuracy": accuracy, "ROC-AUC": roc_auc,
↪ "Train Time": train_time}

def train_simple_mlp(self, X, y, mlp_params=None):
    """
    train simple MLP, playground to test the model with params
    """
    # split up data
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪ 3, random_state=1)

    # create model with params
    mlp = MLPClassifier(**mlp_params)

    # measure training
    start_time = time.time()
    mlp.fit(X_train, y_train)
    train_time = time.time() - start_time

    # eval
    y_pred = mlp.predict(X_test)
    y_prob = mlp.predict_proba(X_test)[: , 1]
    accuracy = accuracy_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_prob)

    return {"Model": "MLP", "Accuracy": accuracy, "ROC-AUC": roc_auc,
↪ "Train Time": train_time}

def train_svm(self, X, y, kernel='rbf'):
    """
    train svm on dataset
    """
    # split data
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪ 3, random_state=1)

    # create modell
    svm = SVC(kernel=kernel, probability=True, random_state=1)

```

```

    # training time
    start_time = time.time()
    svm.fit(X_train, y_train)
    train_time = time.time() - start_time

    # eval
    y_pred = svm.predict(X_test)
    y_prob = svm.predict_proba(X_test)[:, 1]
    accuracy = accuracy_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_prob)

    return {"Model": f"SVM ({kernel})", "Accuracy": accuracy, "ROC-AUC":
    ↪roc_auc, "Train Time": train_time}

def compare_models(self, X, y):
    """
    train and compare the models
    """
    results = []

    # mlp
    mlp_results = self.train_mlp(X, y)
    results.append(mlp_results)

    # svm
    for kernel in self.kernels:
        svm_results = self.train_svm(X, y, kernel)
        results.append(svm_results)

    # create DF
    df_results = pd.DataFrame(results)
    return df_results

def k_fold_cross_validation(self, model_type, X, y, kernel=None, k=5):
    """
    performs k-fold cross validation on X and y data
    """
    if model_type == "MLP":
        model = MLPClassifier(hidden_layer_sizes=self.hidden_layer_sizes,
    ↪max_iter=self.max_iter, random_state=1)
    elif model_type == "SVM":
        model = SVC(kernel=kernel, probability=True, random_state=1)
    else:
        raise ValueError("error bad name")

    # scores

```

```

        scores = cross_val_score(model, X, y, cv=k, scoring='accuracy')
        return scores.mean()

    def evaluate_and_compare(self, dataset_name='moons', noise_levels=[0.1, 0.
↪2, 0.3]):
        """
        eval and compare MLP and SVM models
        """
        results = []

        # iterate over noise levels
        for noise in noise_levels:
            X, y = self.generate_dataset(dataset_name, noise=noise,
↪n_samples=200, random_state=1)

            # train and compare models
            model_results = self.compare_models(X, y)
            model_results['Dataset'] = dataset_name
            model_results['Noise'] = noise
            results.append(model_results)

        # dataframe results
        df_results = pd.concat(results, ignore_index=True)
        return df_results

    def evaluate_and_compare_with_cross_validation(self, dataset_name='moons',
↪noise_levels=[0.1, 0.2, 0.3], k=5):
        """
        evaluate MLP and SVM models with k fold cross validation
        """
        results = []

        # noise levels loop
        for noise in noise_levels:
            X, y = self.generate_dataset(dataset_name, noise=noise,
↪n_samples=200, random_state=1)

            # k Fold CV for MLP
            mlp_cv_score = self.k_fold_cross_validation(model_type="MLP", X=X,
↪y=y, k=k)
            mlp_results = {
                "Model": "MLP",
                "Dataset": dataset_name,
                "Noise": noise,
                "k-Fold CV Accuracy": mlp_cv_score
            }

```

```

        results.append(mlp_results)

        # CV for each SVM kernel
        for kernel in self.kernels:
            svm_cv_score = self.k_fold_cross_validation(model_type="SVM",
↪X=X, y=y, kernel=kernel, k=k)
            svm_results = {
                "Model": f"SVM ({kernel})",
                "Dataset": dataset_name,
                "Noise": noise,
                "k-Fold CV Accuracy": svm_cv_score
            }
            results.append(svm_results)

        # dataframe
        df_results = pd.DataFrame(results)
        return df_results

```

1.0.3 2. Generate and Visualize Datasets

```

[67]: # create the class
      comparison = ModelComparison()

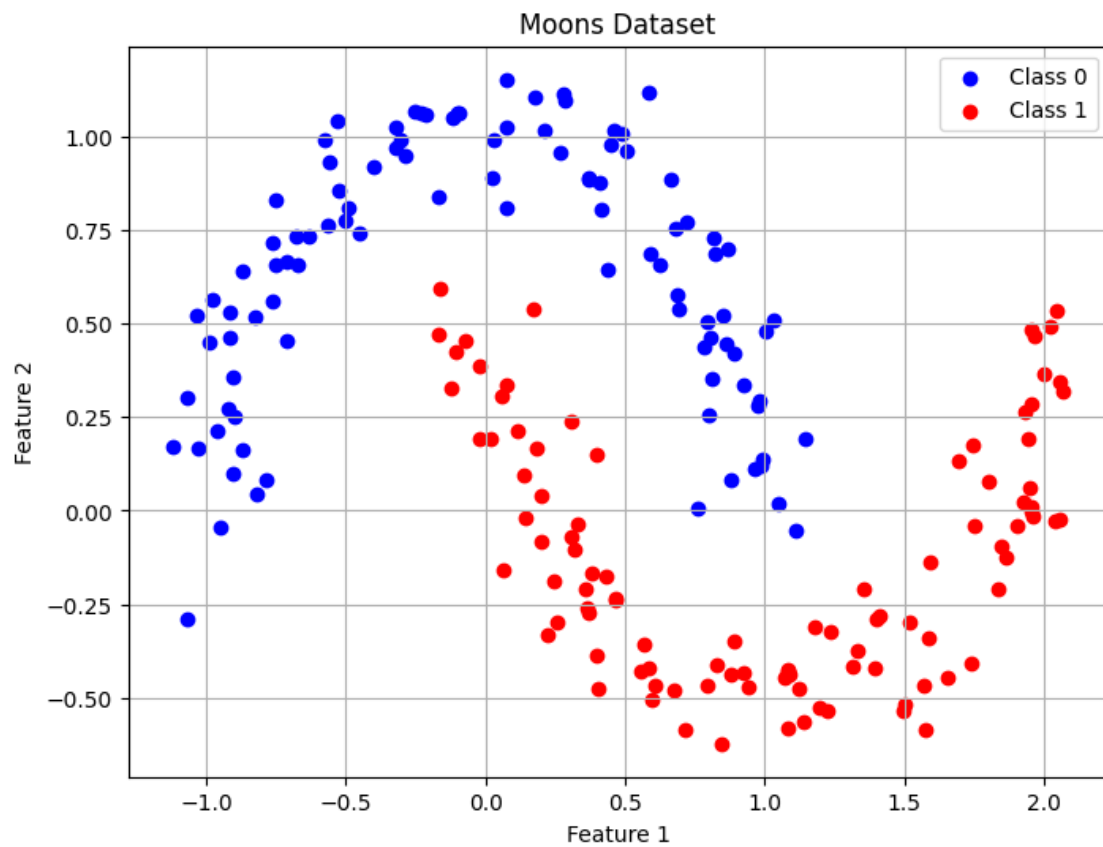
```

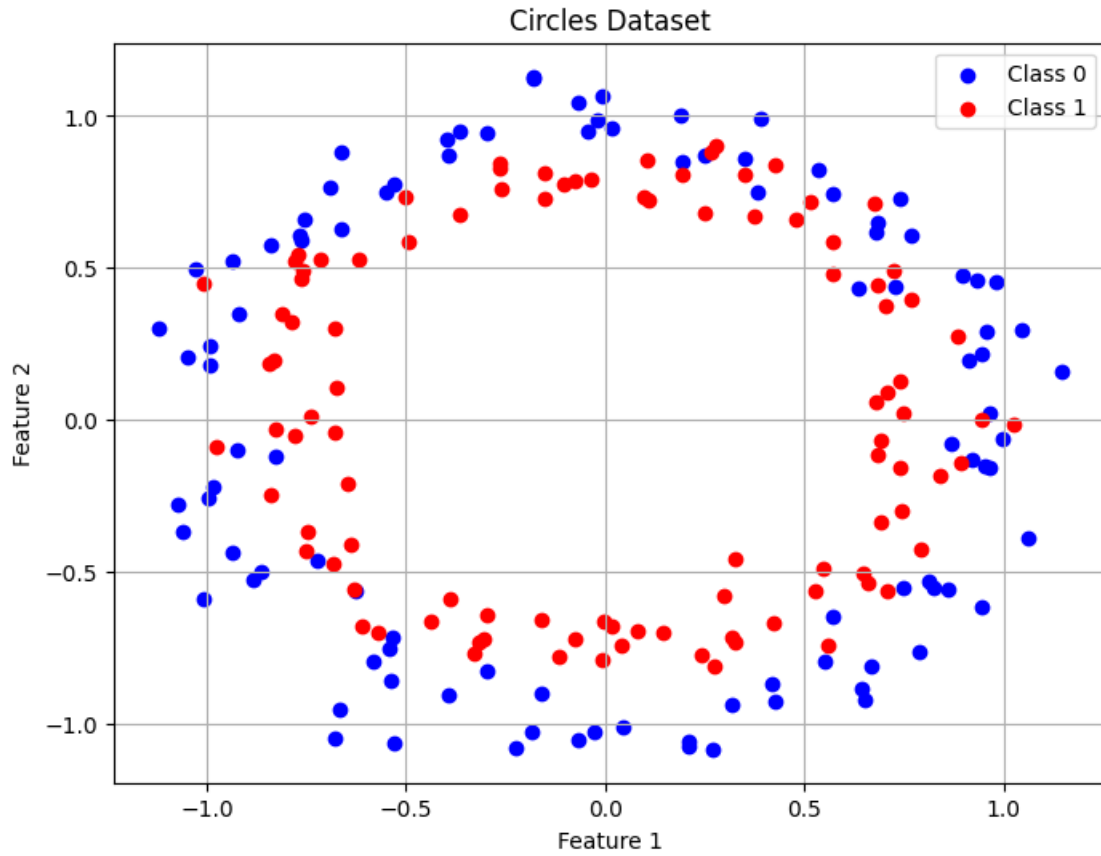
```

[68]: # create the data
      X_moons, y_moons = comparison.generate_dataset('moons', n_samples=200, noise=0.
↪1)
      X_circles, y_circles = comparison.generate_dataset('circles', n_samples=200,
↪noise=0.1)

      # compare the viauslize the results for each dataset
      comparison.visualize_dataset(X_moons, y_moons, "moons")
      comparison.visualize_dataset(X_circles, y_circles, "circles")

```





1.0.4 3. Train Simple MLP

```
[69]: custom_mlp_params = {
        'hidden_layer_sizes': (5,),
        'max_iter': 2000,
        'random_state': 1
    }
    mlp_results = comparison.train_simple_mlp(X_moons, y_moons,
        ↪ mlp_params=custom_mlp_params)
    print("Simple MLP on Moons Dataset:\n", mlp_results)
```

Simple MLP on Moons Dataset:

```
{'Model': 'MLP', 'Accuracy': 0.85, 'ROC-AUC': np.float64(0.9241071428571428),
'Train Time': 0.08712005615234375}
```

```
[71]: custom_mlp_params = {
        'hidden_layer_sizes': (10,),
        'max_iter': 2000,
        'random_state': 1
    }
```



```
mlp_results = comparison.train_simple_mlp(X_moons, y_moons,
    ↪mlp_params=custom_mlp_params)
print("Simple MLP on Moons Dataset:\n", mlp_results)
```

Simple MLP on Moons Dataset:

```
{'Model': 'MLP', 'Accuracy': 0.85, 'ROC-AUC': np.float64(0.9397321428571428),
'Train Time': 0.06336712837219238}
```

```
[72]: custom_mlp_params = {
        'hidden_layer_sizes': (20,),
        'max_iter': 2000,
        'random_state': 1
    }
mlp_results = comparison.train_simple_mlp(X_moons, y_moons,
    ↪mlp_params=custom_mlp_params)
print("Simple MLP on Moons Dataset:\n", mlp_results)
```

Simple MLP on Moons Dataset:

```
{'Model': 'MLP', 'Accuracy': 0.9166666666666666, 'ROC-AUC':
np.float64(0.96875), 'Train Time': 0.08768582344055176}
```

```
[73]: custom_mlp_params = {
        'hidden_layer_sizes': (50,),
        'max_iter': 2000,
        'random_state': 1
    }
mlp_results = comparison.train_simple_mlp(X_moons, y_moons,
    ↪mlp_params=custom_mlp_params)
print("Simple MLP on Moons Dataset:\n", mlp_results)
```

Simple MLP on Moons Dataset:

```
{'Model': 'MLP', 'Accuracy': 0.9666666666666667, 'ROC-AUC':
np.float64(0.9966517857142857), 'Train Time': 0.21885919570922852}
```

```
[74]: custom_mlp_params = {
        'hidden_layer_sizes': (100,),
        'max_iter': 2000,
        'random_state': 1
    }
mlp_results = comparison.train_simple_mlp(X_moons, y_moons,
    ↪mlp_params=custom_mlp_params)
print("Simple MLP on Moons Dataset:\n", mlp_results)
```

Simple MLP on Moons Dataset:

```
{'Model': 'MLP', 'Accuracy': 0.9666666666666667, 'ROC-AUC':
np.float64(0.9988839285714286), 'Train Time': 0.23769426345825195}
```

```
[75]: custom_mlp_params = {
        'hidden_layer_sizes': (200,),
```

```

        'max_iter': 2000,
        'random_state': 1
    }
    mlp_results = comparison.train_simple_mlp(X_moons, y_moons,
        ↪mlp_params=custom_mlp_params)
    print("Simple MLP on Moons Dataset:\n", mlp_results)

```

Simple MLP on Moons Dataset:

```

{'Model': 'MLP', 'Accuracy': 0.9666666666666667, 'ROC-AUC':
np.float64(0.9988839285714286), 'Train Time': 0.23634815216064453}

```

```

[76]: custom_mlp_params = {
        'hidden_layer_sizes': (400,),
        'max_iter': 2000,
        'random_state': 1
    }
    mlp_results = comparison.train_simple_mlp(X_moons, y_moons,
        ↪mlp_params=custom_mlp_params)
    print("Simple MLP on Moons Dataset:\n", mlp_results)

```

Simple MLP on Moons Dataset:

```

{'Model': 'MLP', 'Accuracy': 0.9666666666666667, 'ROC-AUC':
np.float64(0.9988839285714286), 'Train Time': 0.313220739364624}

```

```

[79]: custom_mlp_params = {
        'hidden_layer_sizes': (30,),
        'max_iter': 2000,
        'random_state': 1
    }
    mlp_results = comparison.train_simple_mlp(X_moons, y_moons,
        ↪mlp_params=custom_mlp_params)
    print("Simple MLP on Moons Dataset:\n", mlp_results)

```

Simple MLP on Moons Dataset:

```

{'Model': 'MLP', 'Accuracy': 0.95, 'ROC-AUC': np.float64(0.9966517857142857),
'Train Time': 0.23001503944396973}

```

1.0.5 4. Explain Rationale

- I used single hidden layer with 5-300 neurons to explore. The more neurons, the better the model, however it does plateau at one point. Roughly 50 were sufficient in the MLP model, providing sufficient capacity to learn non-linear patterns in simple datasets like moons and circles while keeping computational cost low.
- Default value of ReLU as the activation function in the hidden layer was efficient for training and to avoid the vanishing gradient problem
- Using the logistic sigmoid activation in the output layer is good for binary classification, as it outputs probabilities. This helped making the model compatible with evaluation metrics like ROC-AUC which was useful for me here.

- Application of k-fold cross validation was used to improve robustness and reliability of the accuracy results. This allowed performance comparison across models and noise levels to select the top candidates.

1.0.6 5. Compare SVM and MLP

```
[80]: # Train and compare models on moons dataset with noise level 0.1
df_moons_results = comparison.evaluate_and_compare(dataset_name='moons')
print("Moons Dataset Results:\n", df_moons_results)
```

Moons Dataset Results:

	Model	Accuracy	ROC-AUC	Train Time	Dataset	Noise
0	MLP	0.883333	0.965208	0.104192	moons	0.1
1	SVM (linear)	0.850000	0.955107	0.006320	moons	0.1
2	SVM (rbf)	1.000000	1.000000	0.000808	moons	0.1
3	MLP	0.900000	0.960718	0.080908	moons	0.2
4	SVM (linear)	0.866667	0.948373	0.001035	moons	0.2
5	SVM (rbf)	0.966667	0.995511	0.000830	moons	0.2
6	MLP	0.883333	0.947250	0.070053	moons	0.3
7	SVM (linear)	0.850000	0.933782	0.001145	moons	0.3
8	SVM (rbf)	0.933333	0.978676	0.000962	moons	0.3

```
[83]: # Train and compare models on circles dataset with noise level 0.1
df_circles_results = comparison.evaluate_and_compare(dataset_name='circles')
print("Circles Dataset Results:\n", df_circles_results)
```

Circles Dataset Results:

	Model	Accuracy	ROC-AUC	Train Time	Dataset	Noise
0	MLP	0.483333	0.561167	0.009141	circles	0.1
1	SVM (linear)	0.450000	0.464646	0.001954	circles	0.1
2	SVM (rbf)	0.833333	0.941639	0.001953	circles	0.1
3	MLP	0.450000	0.547699	0.005161	circles	0.2
4	SVM (linear)	0.450000	0.475870	0.001365	circles	0.2
5	SVM (rbf)	0.700000	0.777778	0.001909	circles	0.2
6	MLP	0.433333	0.524130	0.006775	circles	0.3
7	SVM (linear)	0.450000	0.517396	0.001264	circles	0.3
8	SVM (rbf)	0.616667	0.693603	0.001626	circles	0.3

- Based on the results above, the SVM models train significantly faster than MLP models with RBF SVM achieving training times under 1 millisecond. We can compare this to MLP times range up to 0.1 seconds which is far longer.
- The MLP training time increases with increased noise levels, likely due to the added complexity and iterative weight updates needed
- Interestingly the SVM models are less affected by noise in training speed, showing consistent training times across different noise levels. I had not expected this, but it makes sense.
- The RBF kernel SVM consistently has the lowest computational cost compared to linear SVM and MLP. This is similar to what we saw in task 2.
- Overall, the MLP is more computationally intensive due to its iterative backpropagation process. This makes it less efficient for simple datasets like the ones above.

1.0.7 6. Apply k-Fold cross validation

```
[33]: mlp_cv_score_moons = comparison.k_fold_cross_validation(model_type="MLP",  
    ↪X=X_moons, y=y_moons, k=5)  
print("MLP 5-Fold CV Accuracy (Moons Dataset):", mlp_cv_score_moons)
```

MLP 5-Fold CV Accuracy (Moons Dataset): 0.8800000000000001

```
[34]: svm_rbf_cv_score_moons = comparison.k_fold_cross_validation(model_type="SVM",  
    ↪X=X_moons, y=y_moons, kernel='rbf', k=5)  
print("SVM (RBF) 5-Fold CV Accuracy (Moons Dataset):", svm_rbf_cv_score_moons)
```

SVM (RBF) 5-Fold CV Accuracy (Moons Dataset): 1.0

```
[35]: svm_linear_cv_score_moons = comparison.  
    ↪k_fold_cross_validation(model_type="SVM", X=X_moons, y=y_moons,  
    ↪kernel='linear', k=5)  
print("SVM (Linear) 5-Fold CV Accuracy (Moons Dataset):",  
    ↪svm_linear_cv_score_moons)
```

SVM (Linear) 5-Fold CV Accuracy (Moons Dataset): 0.86

```
[36]: mlp_cv_score_circles = comparison.k_fold_cross_validation(model_type="MLP",  
    ↪X=X_circles, y=y_circles, k=5)  
print("MLP 5-Fold CV Accuracy (Circles Dataset):", mlp_cv_score_circles)
```

MLP 5-Fold CV Accuracy (Circles Dataset): 0.495

```
[37]: svm_rbf_cv_score_circles = comparison.k_fold_cross_validation(model_type="SVM",  
    ↪X=X_circles, y=y_circles, kernel='rbf', k=5)  
print("SVM (RBF) 5-Fold CV Accuracy (Circles Dataset):",  
    ↪svm_rbf_cv_score_circles)
```

SVM (RBF) 5-Fold CV Accuracy (Circles Dataset): 0.835

```
[38]: svm_linear_cv_score_circles = comparison.  
    ↪k_fold_cross_validation(model_type="SVM", X=X_circles, y=y_circles,  
    ↪kernel='linear', k=5)  
print("SVM (Linear) 5-Fold CV Accuracy (Circles Dataset):",  
    ↪svm_linear_cv_score_circles)
```

SVM (Linear) 5-Fold CV Accuracy (Circles Dataset): 0.44000000000000006

```
[39]: df_moons_cv_results = comparison.  
    ↪evaluate_and_compare_with_cross_validation(dataset_name='moons',  
    ↪noise_levels=[0.1, 0.2, 0.3], k=5)  
print("Moons Dataset Cross-Validated Model Comparison:\n", df_moons_cv_results)
```

Moons Dataset Cross-Validated Model Comparison:

	Model	Dataset	Noise	k-Fold CV Accuracy
0	MLP	moons	0.1	0.875

1	SVM (linear)	moons	0.1	0.880
2	SVM (rbf)	moons	0.1	1.000
3	MLP	moons	0.2	0.875
4	SVM (linear)	moons	0.2	0.875
5	SVM (rbf)	moons	0.2	0.955
6	MLP	moons	0.3	0.865
7	SVM (linear)	moons	0.3	0.850
8	SVM (rbf)	moons	0.3	0.910

```
[84]: df_circles_cv_results = comparison.  
      ↪ evaluate_and_compare_with_cross_validation(dataset_name='circles',  
      ↪ noise_levels=[0.1, 0.2, 0.3], k=5)  
      print("Circles Dataset Cross-Validated Model Comparison:\n",  
      ↪ df_circles_cv_results)
```

Circles Dataset Cross-Validated Model Comparison:

	Model	Dataset	Noise	k-Fold CV Accuracy
0	MLP	circles	0.1	0.480
1	SVM (linear)	circles	0.1	0.385
2	SVM (rbf)	circles	0.1	0.830
3	MLP	circles	0.2	0.475
4	SVM (linear)	circles	0.2	0.400
5	SVM (rbf)	circles	0.2	0.635
6	MLP	circles	0.3	0.445
7	SVM (linear)	circles	0.3	0.415
8	SVM (rbf)	circles	0.3	0.560

1.0.8 7. Conclusion

- In this project I designed a minimal MLP with a single hidden layer of 5-300 neurons compare the impact and to efficiently capture non-linear patterns in the moons and circles datasets.
- To maintain the requirement of “simple”, I used the standard ReLU activation in the hidden layer which is a sigmoid output for binary classification. In addition, I used the Adam optimizer that the MLP was designed to learn quickly and avoid vanishing gradient issues.
- Overall SVM models particularly with the RBF kernel showed significantly faster training times than MLP across all noise levels, making SVM more computationally efficient for these tasks. We can also see this with the times that were captured above.
- k-fold cross-validation was applied for both MLP and SVM models to get more reliable accuracy estimates.
- Overall while SVM models are faster and computationally more efficient, the simple MLP provides flexibility for handling increased noise.